

ECE 590-03: Homework #1

Linear Model, Back Propagation and Building a CNN

Hai Li

ECE Department, Duke University — August 23, 2020

Objectives

Homework #1 covers the contents of Lectures 02 ~ 04. This assignment includes conceptual questions on the linear model, back propagation and convolutions, as well as lab questions involving trying out LMS algorithm and building and observing a CNN model.

The lab questions in this assignment don't involve the training of deep learning models. So you are welcome to use your own computer to finish the labs. Please refer to the [NumPy/PyTorch tutorial](#) slides on Sakai for the environment setup on your computer. If you intend to use the OIT server, please refer to the appendix on Page 6 of this assignment document for the environment setup.



Warning: You are asked to complete the assignment independently.

This lab has a total of 100 points plus 10 bonus points, yet your final score cannot exceed 100 points. You must submit your report in PDF format and your original codes for the lab questions through Sakai before **11:55:00pm, Thursday, September 3**. You need to submit **three individual files** including (1) *a self-contained report in PDF format* that provides answers to all the conceptual questions and clearly demonstrates all your lab results and observations, (2) *a single code file* used to produce all the results for Lab: LMS algorithms; and (3) *a Jupyter notebook file* for Lab: LeNet-5. **Note that 20 percent of the grade will be deducted for the submissions uploaded in a zip file.**

1 True/False Questions (10 pts)

For each question, please provide a short explanation to support your judgment.

Problem 1.1 (2 pts) Using a proper setting (e.g., learning rate, initialization etc.), the gradient descent algorithm is always able to yield an optimal solution on convex functions, as the local optimal value is always the global optimal value.

Problem 1.2 (2 pts) For a given task, having more layers in a neural network model always improves its capacity. Thus, it is beneficial to have as many layers as possible to improve NN's generalization ability.

Problem 1.3 (2 pts) Given a learning task that can be perfectly learned by a Madaline model, the same set of weight values will be achieved after training, no matter how the Madaline is initialized.

Problem 1.4 (2 pts) Sigmoid neurons are likely to 'die' during the training process, as the gradients in the middle of the 'S-shape' sigmoidal curve may vanish. Thus, these neurons are less likely to be updated. We should use a larger learning rate to train all of the sigmoid neurons.

Problem 1.5 (2 pts) According to the "convolution shape rule," for a convolution operation with fixed kernel size, stride and padding, increasing the height and width of the input feature map will always lead to a larger output feature map size.

2 Adalines (15 pts)

In the following problems, you will be asked to derive the output of a given Adaline, or propose proper weight values for the Adaline to mimic the functionality of some simple logic functions. For all problems, please consider +1 as **True** and -1 as **False** in the inputs and outputs.

Problem 2.1 (3 pts) Observe the Adaline shown in Figure 1, fill in the feature s and output y for each pair of inputs given in the truth table. What logic function is this Adaline performing?




Figure 1: Problem 2.1.

Problem 2.2 (4 pts) Propose proper values for weight w_0, w_1 and w_2 in the Adaline shown in Figure 2 to perform the functionality of a logic **NOR** function. Fill in the feature s for each pair of inputs given in the truth table to prove the functionality is correct.




Figure 2: Problem 2.2.

Problem 2.3 (4 pts) Propose proper values for weight w_0, w_1, w_2 and w_3 in the Adaline shown in Figure 3 to perform the functionality of a **Majority Vote** function. Fill in the feature s for each triplet of inputs given in the truth table to prove the functionality is correct.

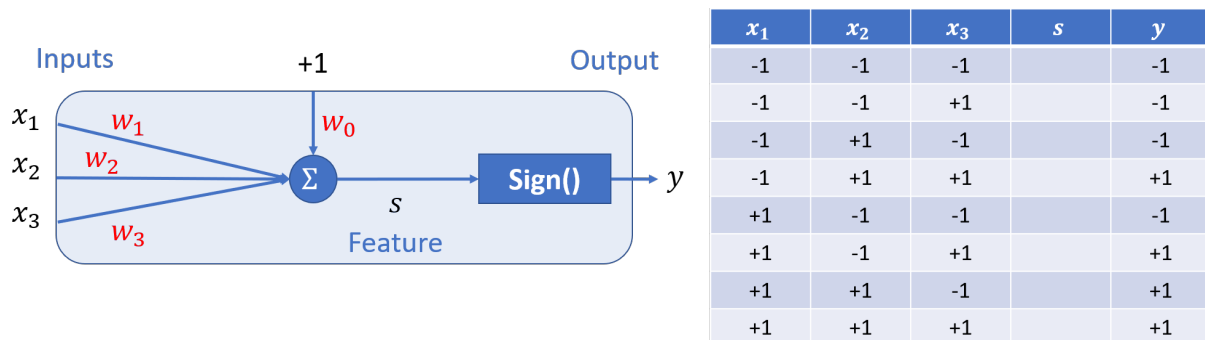


Figure 3: Problem 2.3.

Problem 2.4 (4 pts) As discussed in Lecture 2, the XOR function cannot be represented with a single Adaline, but can be represented with a 2-layer Madaline. Propose proper values for second-layer weight w_{20}, w_{21} and w_{22} in the Madaline shown in Figure 4 to perform the functionality of a **XOR** function. Fill in the feature s for each pair of inputs given in the truth table to prove the functionality is correct.

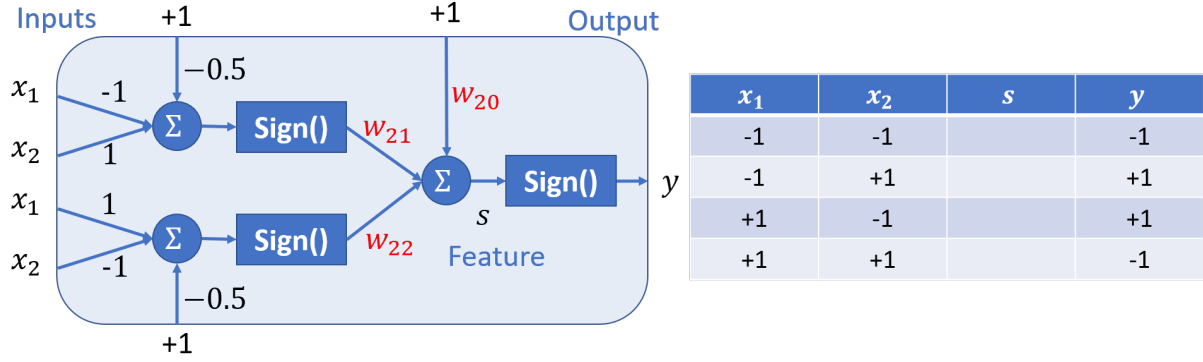


Figure 4: Problem 2.4.

3 Back Propagation (10 pts)

Problem 3.1 (5 pts) The feed-forward computation of a logistic neuron could be represented as follows. Given an input $x_l \in \mathcal{R}^{n \times 1}$, the output $x_{l+1} \in \mathcal{R}^{m \times 1}$ is calculated as follows:

$$y_l = W_l x_l + b_l, \quad x_{l+1} = \sigma(y_l). \quad (1)$$

Where $W \in \mathcal{R}^{m \times n}$ denotes the weight matrix, $b \in \mathcal{R}^{m \times 1}$ denotes the bias term, and $\sigma : \mathcal{R}^{m \times 1} \Rightarrow \mathcal{R}^{m \times 1}$ denotes the Sigmoid activation function. Following the chain rule, derive the gradient $\frac{\partial x_{l+1}}{\partial x_l}, \frac{\partial x_{l+1}}{\partial W_l}, \frac{\partial x_{l+1}}{\partial b_l}$ in a **vectorized format**¹. (Hint: make sure the gradient has the same shape as the parameter.)

Problem 3.2 (5 pts) Consider a 2-layer fully-connected NN composed by the logistic neuron mentioned in Problem 3.1, where we have input $x_1 \in \mathcal{R}^{n \times 1}$, hidden feature $x_2 \in \mathcal{R}^{m \times 1}$, output $x_3 \in \mathcal{R}^{k \times 1}$ and weights and bias W_1, W_2, b_1, b_2 of the two layers. A MSE loss function $L = \frac{1}{2}(t - x_3)^T(t - x_3)$ is applied in the end. Following the chain rule, derive the gradient $\frac{\partial L}{\partial W_1}$ in a **vectorized format**.

4 2D Convolution (10 pts)

Problem 4.1 (5 pts) Derive the 2D convolution results of the following 5×9 input matrix and the 3×3 kernel. Consider 0s are padded around the input and the stride is 1, so that the output should also have shape 5×9 .

$$\begin{bmatrix} 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 1 & 1 & 1 & 1 & 0 & 0 \\ 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ 0 & 0 & 1 & 1 & 1 & 1 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 \end{bmatrix} * \begin{bmatrix} 0 & -1/4 & 0 \\ -1/4 & 2 & -1/4 \\ 0 & -1/4 & 0 \end{bmatrix}$$

Problem 4.2 (5 pts) Compare the output matrix and the input matrix in Problem 4.1, briefly analyze the effect of this 3×3 kernel on the input.

$$\text{Hint: } \begin{bmatrix} 0 & -1/4 & 0 \\ -1/4 & 2 & -1/4 \\ 0 & -1/4 & 0 \end{bmatrix} = \begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix} + \left(\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix} - \begin{bmatrix} 0 & 1/4 & 0 \\ 1/4 & 0 & 1/4 \\ 0 & 1/4 & 0 \end{bmatrix} \right)$$

5 Lab: LMS Algorithm (15 pts)

In this lab question, you will implement the LMS algorithm with NumPy to learn a linear regression model for the provided dataset. You will also be directed to analyze how the choice of learning rate in the LMS

¹Please refer to <http://cs231n.stanford.edu/handouts/linear-backprop.pdf>, if you are not familiar with deriving gradients in vectorized form.

algorithm affect the final result. All the codes generating the results of this lab should be gathered in one file and submit to Sakai.

Lab 1 (15 pts)

To start with, please download the `dataset.mat` file from Sakai and load it into NumPy arrays^a. There are two variables in the file: data $X \in \mathbb{R}^{100 \times 3}$ and target $D \in \mathbb{R}^{100 \times 1}$. Each individual pair of data and target is composed into X and D following the same way as discussed on Lecture 2 Page 8. Specifically, each row in X correspond to the transpose of a data point, with the first element as constant 1 and the other two as the two input features x_{1k} and x_{2k} . The goal of the learning task is finding the weight vector $W \in \mathbb{R}^{3 \times 1}$ for the linear model that can minimize the MSE loss, which is also formulated on Lecture 2 Page 8.

- (a) (3pt) Directly compute the least square (Wiener) solution with the provided dataset. What is the optimal weight W^* ? What is the MSE loss of the whole dataset when the weight is set to W^* ?
- (b) (4pt) Now consider that you can only train with 1 pair of data point and target each time. In such case, the LMS algorithm should be used to find the optimal weight. Please initialize the weight vector as $W^0 = [0, 0, 0]^T$, and update the weight with the LMS algorithm. After each *epoch* (every time you go through all the training data and loop back to the beginning), compute and record the MSE loss of the current weight on the whole dataset. Run LMS for 20 epochs with learning rate $r = 0.01$, report the weight you get in the end and plot the MSE loss in log scale vs. Epochs.
- (c) (3pt) Scatter plot the points (x_{1k}, x_{2k}, d_k) for all 100 data-target pairs in a 3D figure^b, and plot the lines corresponding to the linear models you got in (a) and (b) respectively in the same figure. Observe if the linear models fit the data well.
- (d) (5pt) Learning rate r is an important hyperparameter for the LMS algorithm, as well as for CNN optimization. Here, try repeat the process in (b) with r set to 0.005, 0.05 and 0.5 respectively. Together with the result you got in (b), plot the MSE losses of the 4 sets of experiments in log scale vs. Epochs in one figure. Then try further enlarge the learning rate to $r = 1$ and observe how the MSE changes. Base on these observations, comment on how learning rate affects the speed and quality of the learning process. (Note: The learning rate tuning for the CNN optimization will be introduced in Lecture 7.)

^aYou may refer to <https://docs.scipy.org/doc/scipy/reference/generated/scipy.io.loadmat.html> for loading matrices in .mat file into NumPy arrays.

^bPlease refer to <https://jakevdp.github.io/PythonDataScienceHandbook/04.12-three-dimensional-plotting.html> for plotting 3D plots with Matplotlib.

6 Lab: LeNet-5 (40 pts)

LeNet-5 is a classic convolutional neural network (CNN) for hand-written digits recognition. It has 2 CONV layers, 2 POOL layers and 3 FC layers. Here we consider an adapted LeNet-5 model designed to classify CIFAR-10 images with size $32 \times 32 \times 3$. Details of the model architecture is given in Table 1. This lab question focuses on building the model in PyTorch and observing the shape of each layer's input, weight and output. Please refer to the **NumPy/PyTorch Tutorial slides** on Sakai and the official documentations if you are unfamiliar with PyTorch syntax. You will have your hands on training this model in Homework #2.

Please finish this lab by completing the `LeNet5.ipynb` notebook file provided on Sakai. The completed notebook file should be submitted to Sakai.

Name	Type	Kernel size	depth/units	Activation	Strides
Conv 1	Convolution	5	6	ReLU	1
MaxPool	MaxPool	2	N/A	N/A	2
Conv 2	Convolution	5	16	ReLU	1
MaxPool	MaxPool	2	N/A	N/A	2
FC1	Fully-connected	N/A	120	ReLU	N/A
FC2	Fully-connected	N/A	84	ReLU	N/A
FC3	Fully-connected	N/A	10	None	N/A

Table 1: Adapted LeNet-5 Model. No padding is applied on both convolution layers. A flatten layer is required before FC1 to reshape the feature.

Lab 2 (50 points)

In the notebook, first run through the first two code blocks, then follow the instructions in the following questions to complete each code block and acquire the answers.

- (a) (10pt) Complete code block 3 for defining the adapted LeNet-5 model. Note that customized CONV and FC classes are provided in code block 2 to replace the `nn.Conv2d` and `nn.Linear` classes in PyTorch respectively. The usage of the customized classes are exactly the same as their PyTorch counterparts, the only difference is that in the customized class the input and output feature maps of the layer will be stored in `self.input` and `self.output` respectively after the forward pass, which will be helpful in question (b). After the code is completed, run through the block and make sure the model forward pass in the end throw no errors. Please copy your code of the completed LeNet5 class into the report PDF.
- (b) (30pt) Complete the for-loop in code block 4 to print the shape of the input feature map, output feature map and the weight tensor of the 5 convolutional and fully-connected layers when processing a single input. Then compute the number of parameters and the number of MACs in each layer with the shapes you get. In your report, use your results to fill in the blanks in Table 2.

Layer	Input shape	Output shape	Weight shape	# Param	# MAC
Conv 1					
Conv 2					
FC1					
FC2					
FC3					

Table 2: Results of Lab 2(b).

Lab 3 (Bonus 10 points)

Please first finish all the required codes in Lab 2, then proceed to code block 5 of the notebook file.

- (a) (2pt) Complete the for-loop in code block 5 to plot the histogram of weight elements in each one of the 5 convolutional and fully-connected layers.
- (b) (3pt) In code block 6, complete the code for backward pass, then complete the for-loop to plot the histogram of weight elements' gradients in each one of the 5 convolutional and fully-connected layers.
- (c) (5pt) In code block 7, finish the code to set all the weights to 0. Perform forward and backward pass again to get the gradients, and plot the histogram of weight elements' gradients in each one of the 5 convolutional and fully-connected layers. Comparing with the histograms you got in (b), are there any differences? Briefly analyze the cause of the difference, and comment on how will initializing CNN model with zero weights will affect the training process. (Note: The CNN initialization methods will be introduced in Lecture 6.)

Appendix: Using the OIT Server

If you wish to finish the Lab questions on the OIT server, please visit <https://vm-manage.oit.duke.edu/containers> and log into your Jupyter Notebook Environment. You can upload the files you need to the server by clicking the button shown in Figure 5:

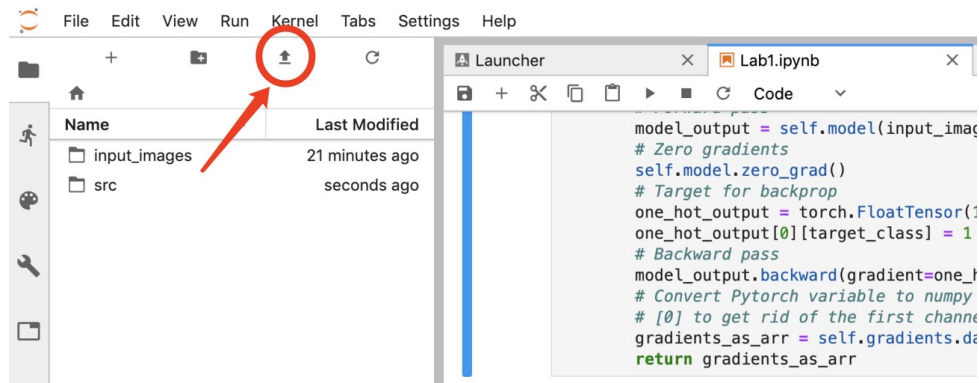


Figure 5: Uploading files Instruction

If you are uploading an zip file, you may unzip it on the server by:

- Press the '+' button and click on "terminal" in the right-hand side "Launcher" column.
- In the terminal, type `unzip *.zip`



Notice: After finishing the lab, please make sure you kill your current process by right-clicking on the `.ipynb` file and select "Shutdown Kernel", as shown in Figure 6:

Please note that there is a 30-minute idle timeout for GPU access set on the OIT server. If you find that you can no longer access the GPU due to the timeout, simply save your progress, log out, restart your browser and log back in, then you can keep working again.

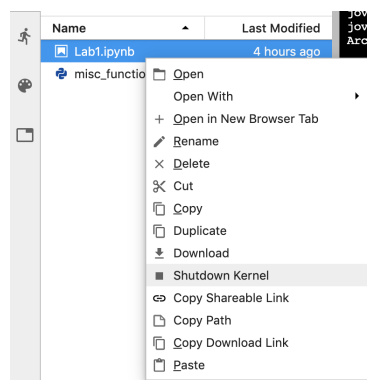


Figure 6: Shutdown kernel before exiting